Discussion 4

Alejandro D. Osborne

July 4, 2019

Mitigating the Harm of Recommender Systems

Recommender Systems are based on user-driven behavior and content, thus these systems are vulnerable to malicious users with dubious intent. The harm caused by Recommender Systems is rather prevalent in today's society. The NYTimes article given talks about how YouTube's algorithm can lead one to videos produced by ideologists that exist on the fringe of the political spectrum - exposing users to ever-more-sensational programming. As an example of the damage, the author provided a link to a study titled Selective Exposure to Misinformation: Evidence from the consumption of fake news during the 2016 U.S. presidential campaign that concluded that misinformation favored Trump more than Clinton in the 2016 election.

According to this study, the top 10 apparent promulgators of misinformation ("fake news" articles) are big-time websites we all know and use - Google, YouTube, and gmail among them. The possible harm caused here is difficult to quantify but it's clear that harm was the outcome.

While there is no easy solution to this problem, data would need to be gathered to identify when traffic is being directed to sites known for spreading misinformation and label the results accordingly. Human oversight could be effective but it's an unlikely option for big tech. For each of these suggestions there are almost as many reasons why they'd be huge problems (e.g. criteria for labeling a "harmful site" could also be hacked, human intervention would be expensive/unsustainable).

Making truth as exciting and profitable as misinformation is already an impossible task without algorithms but with them it makes this problem near impossible to rectify.